

Owl: Performance-Aware Scheduling for Resource-Efficient Function-as-a-Service Cloud

Huangshi Tian¹, Suyi Li¹, Ao Wang^{2,3}, Wei Wang¹, Tianlong Wu³, Haoran Yang³

¹HKUST, ²George Mason University, ³Alibaba Cloud

FaaS Gaining Popularity

“

A new report from Datadog has found that serverless computing could be entering the **mainstream** with **over half** of all organizations using serverless ...

—— TechCrunch¹, Jun. 2022

”

“

AWS Lambda ..., and more than **a million** customers are using it today, according to AWS.

—— Protocol², Aug. 2022

”

1. [Datadog finds serverless computing is going mainstream](https://tcn.ch/3D5GhHB), <https://tcn.ch/3D5GhHB>

2. [Amazon's Werner Vogels: Enterprises are more daring than you might think](https://bit.ly/3F8Xtij), <https://bit.ly/3F8Xtij>

FaaS Gaining Popularity

“

A new report from Datadog has found that serverless computing could be entering the **mainstream** with **over half** of all organizations using

Problem: How to serve functions efficiently?

”

“

AWS Lambda ..., and more than **a million** customers are using it today, according to AWS.

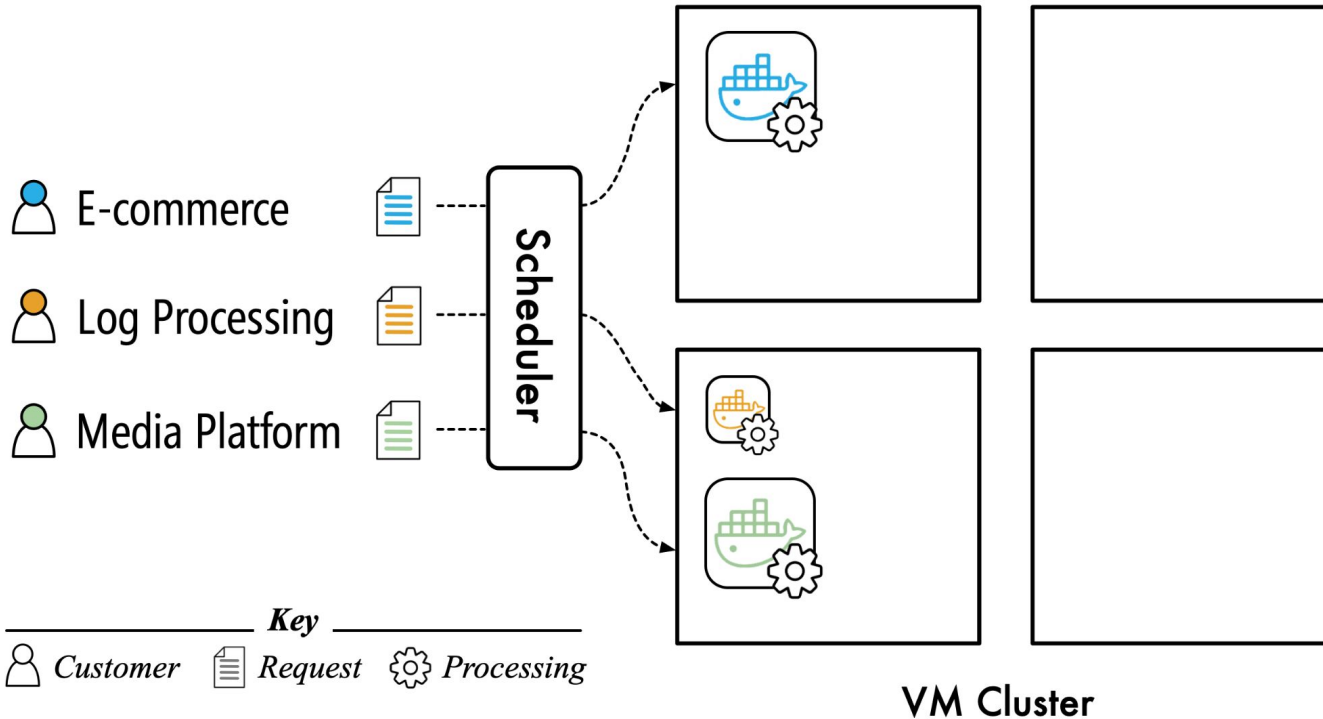
—— Protocol², Aug. 2022

”

1. Datadog finds serverless computing is going mainstream, <https://tcn.ch/3D5GhHB>

2. Amazon's Werner Vogels: Enterprises are more daring than you might think, <https://bit.ly/3F8Xtij>

Scheduling in Public FaaS



Scheduling in Public FaaS

FaaS Scheduling

||


Request Routing

which sandbox to serve request?

+

Sandbox Placement

which VM to place sandbox?



Scheduling in Public FaaS

FaaS Scheduling

||


Request Routing

which sandbox to serve request?

+

Sandbox Placement

which VM to place sandbox?



Scheduling in Public FaaS

FaaS Scheduling

||

Request Routing

which sandbox to serve request?

+

Sandbox Placement

which VM to place sandbox?

Model*

1. Each **VM** has a memory capacity.
2. Each **sandbox** has a memory size.

Goal

Pack sandboxes onto VMs.

* similar as bin-packing

Status Quo

FaaS Scheduling

||

Request Routing

which sandbox to serve request?

+

Sandbox Placement

which VM to place sandbox?

Setting*

1. Each **VM** has a memory capacity.
2. Each **sandbox** has a memory size.

Goal

Pack sandboxes onto VMs.

* similar as bin-packing

State of the Practice

||

Most-Recently Used


for keeping more sandboxes idle

+

First Fit

for quick decision

Resource Inefficiency



* data collected from an one-day production trace

State of the Practice

||

Most-Recently Used


for keeping more sandboxes idle

+

First Fit

for quick decision


Naive Overcommitment



Normal VM


—**Sandbox Overcommitment**→

usage-based heuristic:
allocate what sandbox utilizes




Overcommitted VM

Naive Overcommitment Falls Short




Normal VM




P95 Lat.¹

—**Sandbox Overcommitment**→

usage-based heuristic:
allocate what sandbox utilizes



Ovct.'ed VM




Ovct.'ed P95 Lat.¹

Performance Degrades


unacceptable in production

1. The latency is measured from a benchmark workload. (See §7.1 in the paper).

Naive Overcommitment Falls Short




—Sandbox Overcommitment→
usage-based heuristic:



Objective:
Resource-Efficient and Performance-Aware FaaS Scheduling

Normal VM



P95 Lat.¹

Performance Degrades

unacceptable in production

Overcommitted VM



Overcommitted P95 Lat.¹


1. The latency is measured from a benchmark workload. (See §7.1 in the paper).

Outline


1. Background and Motivation
2. Collocation Profiling
3. Profile-Guided Overcommitment
4. Performance-Monitored Overcommitment
5. Owl and its Evaluation

Restrictions for Profiling

No latency **degradation**.




No **offline** invocation (b/c of side-effect).




Restrictions for Profiling

No latency **degradation**.



No **offline** invocation (b/c of side-effect).



Ideal profiling:



in-production



performance-protected

New Technique: Collocation Profiling

Key Question

How many sandboxes can a VM host?

☐ in-production

☐ performance-protected

Collocation Profiling

Key Question

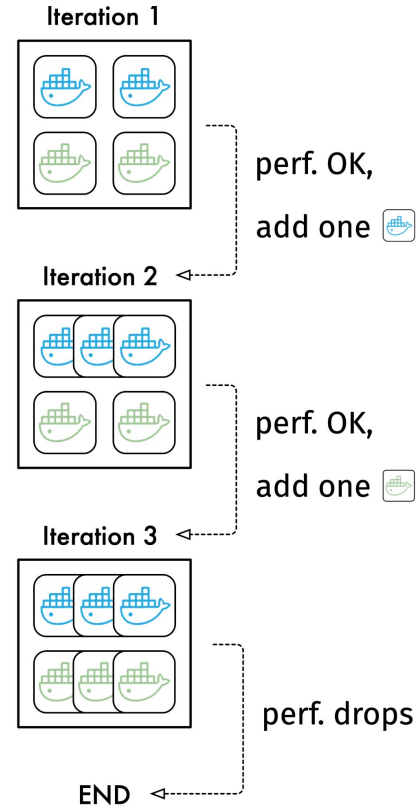
How many sandboxes can a VM host?

Procedure

1. **Saturate** sandboxes with requests.
2. Iteratively add more sandboxes ...
3. ... until perf. starts dropping.

☐ in-production

☐ performance-protected



Collocation Profiling

Key Question

How many sandboxes can a VM host?

Procedure

1. **Saturate** sandboxes with requests.
2. Iteratively add more sandboxes ...
3. ... until perf. starts dropping.



in-production

cf. Step 1, 2



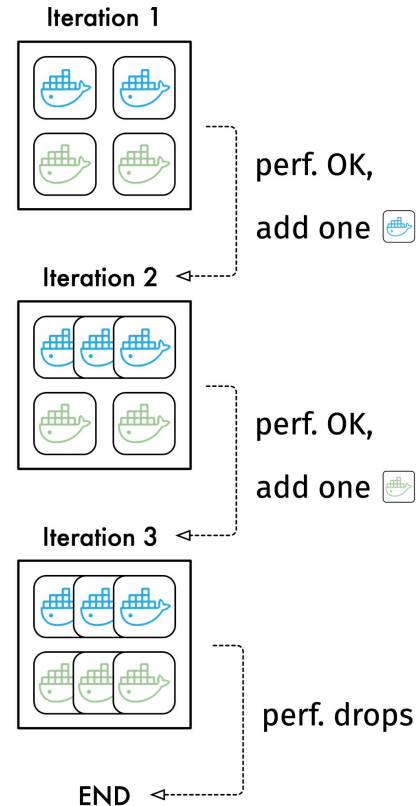
performance-protected

cf. Step 3












resource-overcommitted

cf. Step 2

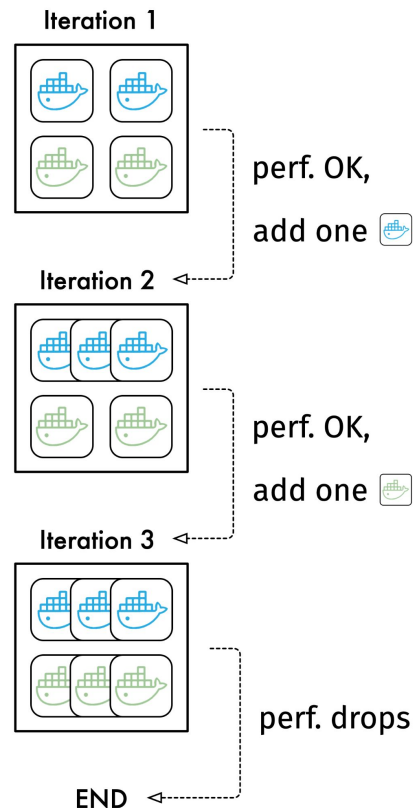


Collocation Profiling










Collocation ²	Util. ¹
 x3  x3	1.48
 x3  x5	1.42
 x3  x3	1.26
 x5	1.22
 x5	1.24
 x7	1.20


1. Util. = Allocated Memory / Total VM Memory
(> 1 means *overcommitment*)

2. Limited to two functions b/c of complexity.



Profile-Guided Overcommitment

Collocation	Util.
 x3  x3	1.48
 x3  x5	1.42
 x3  x3	1.26
 x5	1.22
 x5	1.24
 x7	1.20












VM Cluster

Key Question

How to place sandboxes using profiles?

Profile-Guided Overcommitment

Collocation	Util.
 x3  x3	1.48
 x3  x5	1.42
 x3  x3	1.26
 x5	1.22
 x5	1.24
 x7	1.20










Greedy Algorithm

Collocating sandboxes with highest util.




VM Cluster

Profile-Guided Overcommitment

Collocation	Util.
 x3  x3	1.48
 x3  x5	1.42
 x3  x3	1.26
 x5	1.22
 x5	1.24
 x7	1.20

Greedy Algorithm

Collocating sandboxes with highest util.




VM Cluster

Profile-Guided Overcommitment

Offline

Collocating sandboxes with highest util.



VM Cluster

Profile-Guided Overcommitment


Offline

Collocating sandboxes with highest util.

Online

1. Periodically update placement.
2. Only include VMs with **sandbox change**.

[more details and optimizations in paper]




VM Cluster

Profile-Guided Overcommitment

Problem

profiling requires continuous requests

⇒ only applies to popular functions




VM Cluster

Performance-Monitored Overcommitment

Solution

1. Usage-based overcommitment.
2. Keep **monitoring** performance.
3. **Remedy** degradation (e.g., sandbox migration).





VM Cluster

Problem: External Degradation

Solution

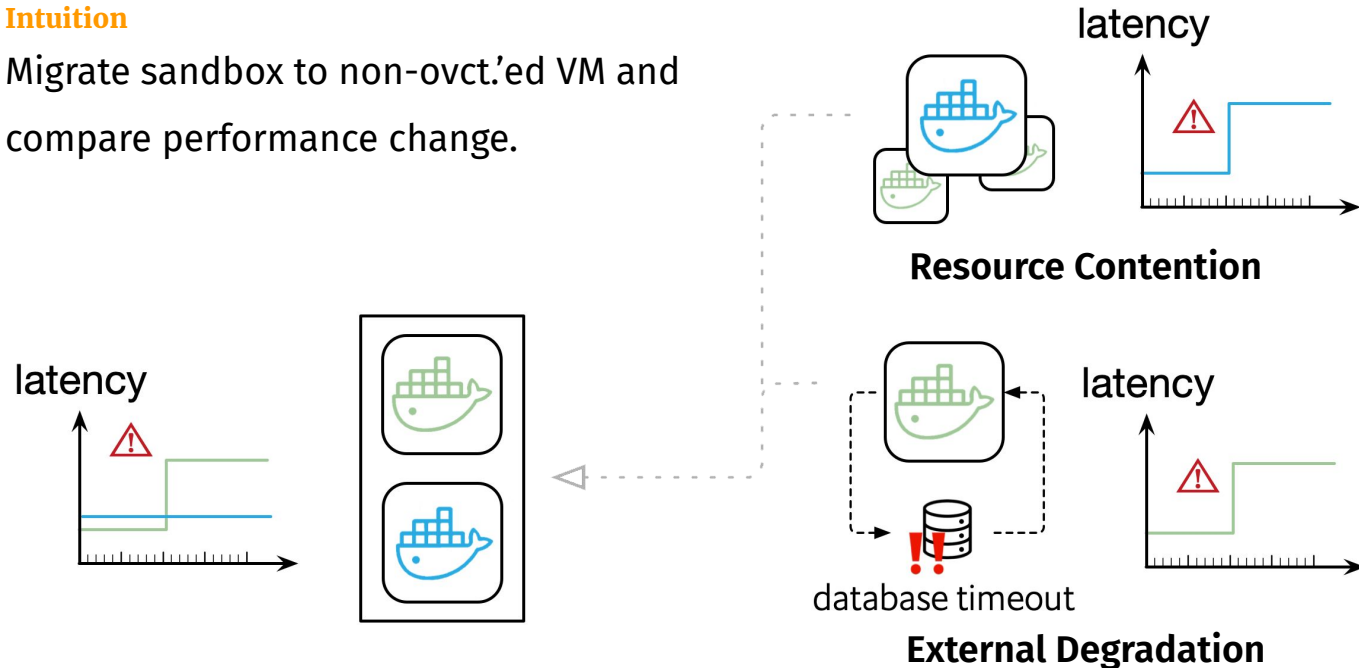
1. Usage-based overcommitment.
2. Keep **monitoring** performance.
3. **Remedy** degradation (e.g., sandbox migration).



New Technique: Comparative Validation

Intuition

Migrate sandbox to non-ovct.'ed VM and compare performance change.




Putting it all Together

Popular Functions

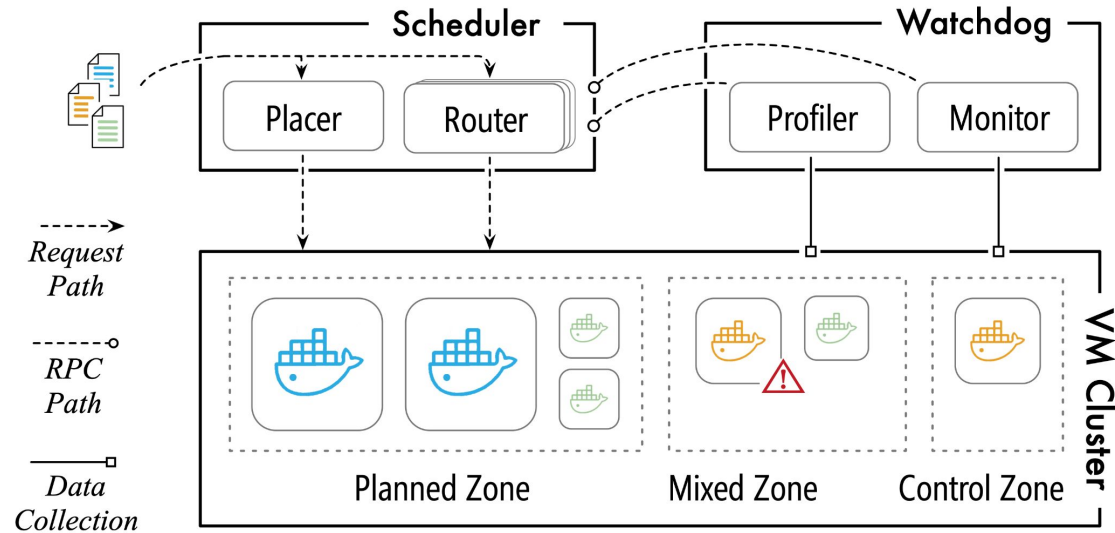
1. Profile Collocations
2. Collocate Sandboxes
3. Consolidate Idle Ones

Less Popular Functions

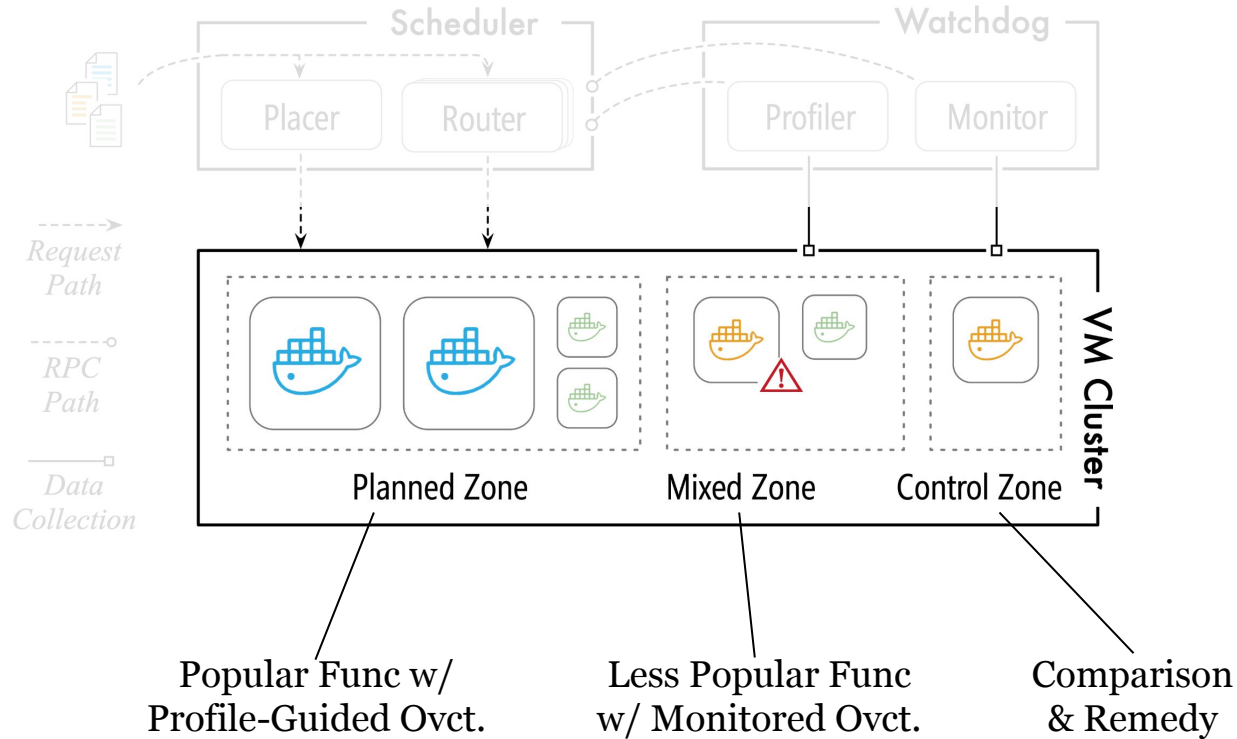
1. Monitor Performance
2. Remedy Degradation
3. Validate Cause



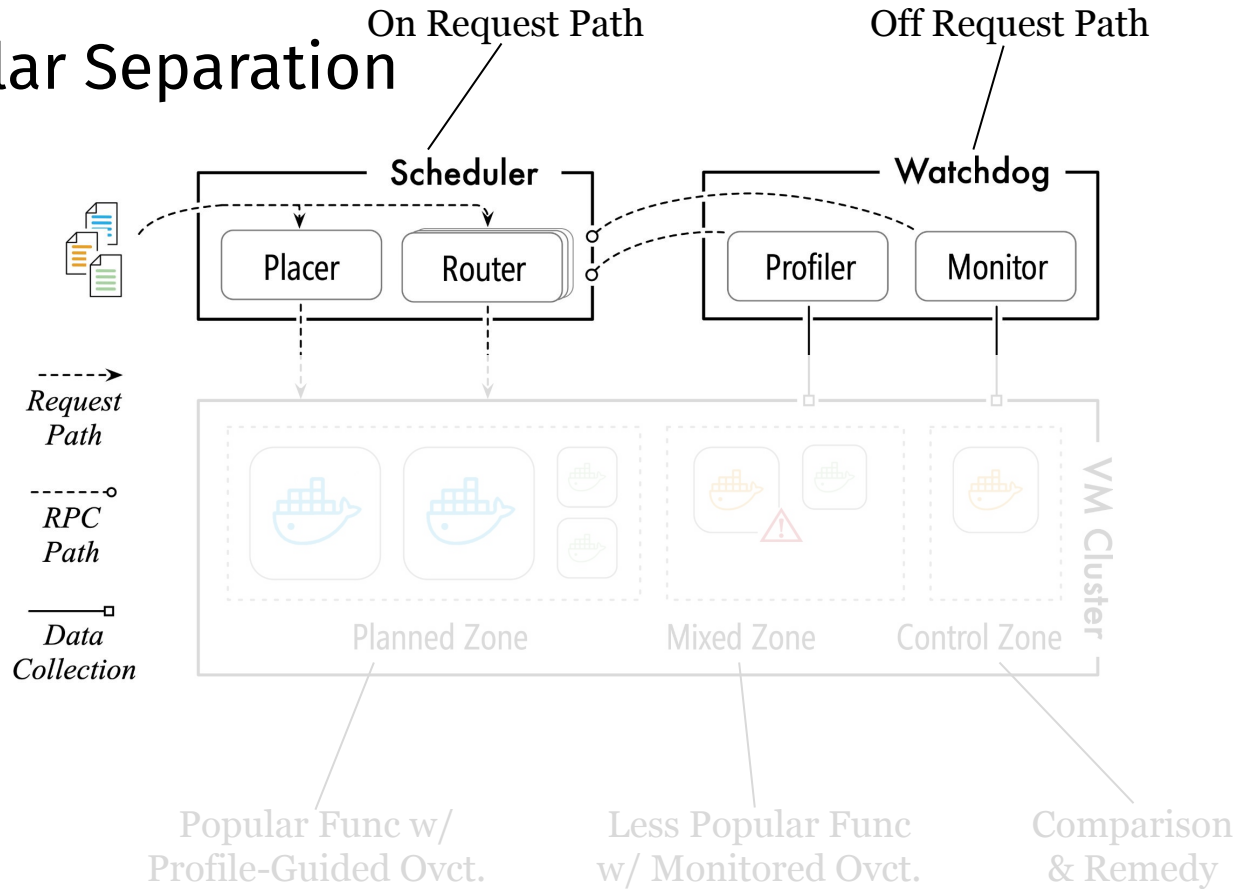
Implementation



Cluster Zoning



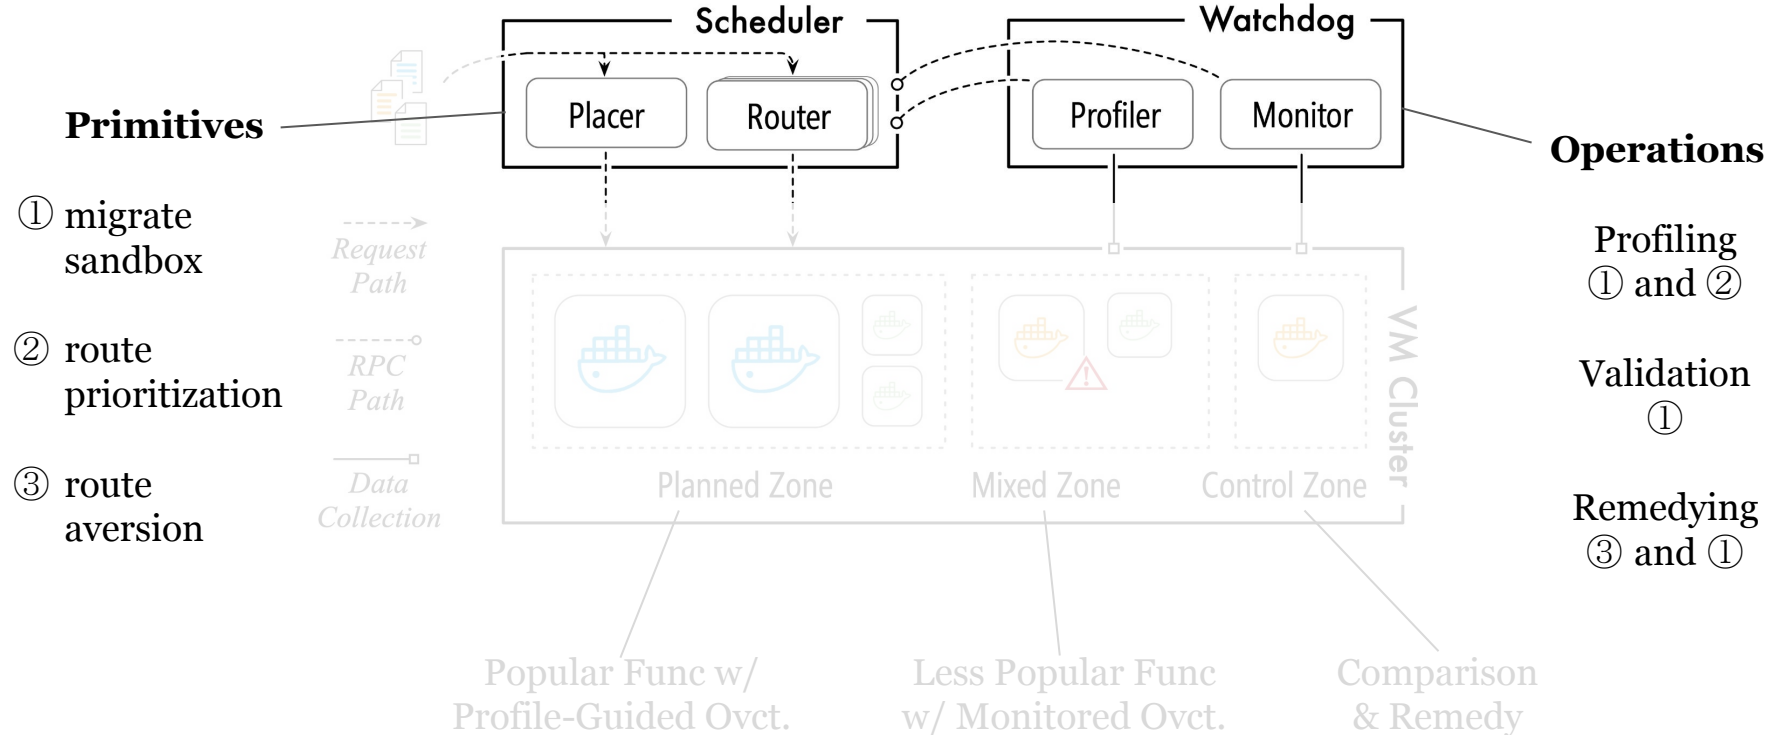
Modular Separation



Primitive-Based Design

On Request Path

Off Request Path



Evaluation Highlight

Setting

- three 20-minute production workloads
- four variant schedulers with different techniques enabled

More results in the paper.

- (1) scheduling latency
- (2) microbenchmarking
- (3) large-scale validation

Results

VM Cost

Function Performance

① Baseline

-

not degraded

② Naive Overcommitment

decrease by 44.26% from ①

degraded

③ ... /w Profile Guidance

increase by 9.24% from ②

not degraded

④ ... /w Consolidation

decrease by 8.53% from ③

not degraded

Overall

decrease by 43.8% from ①

not degraded

Open-Source Benchmark

Abbreviation	Function	Memory Size	Actual Usage	Lanaguage	Dependencies
QV	Query Vacancy	256 MiB	~70MiB	JavaScript	Key-Value Store
RS	Reserve Spot	256 MiB	~70MiB	JavaScript	Key-Value Store, Message Queue
AL	Anonymize Log	1024 MiB	~20MiB	Rust	Message Queue
FL	Filter Log	1024 MiB	~20MiB	Rust	Message Queue
DO	Detect Object	3072 MiB	~1700MiB	Python	Model Serving Framework
CI	Classify Image	2560 MiB	~500MiB	Python	Model Serving Framework
GMM	Get Media Meta	128 MiB	~20MiB	Python	Object Store
CA	Convert Audio	256 MiB	~100MiB	Python	Object Store
ID	Ingest Data	768 MiB	~10MiB	C++	SQL Database
DA	Detect Anomaly	768 MiB	~10MiB	C++	SQL Database

<https://github.com/All-less/faas-scheduling-benchmark>

Conclusion

Performance-Aware and Resource-Efficient Scheduling in Public FaaS

1. Profile-Guided Overcommitment
2. Performance-Monitored Overcommitment